

# NLP HW1 Report

In order to complete the assignment I have followed the following actions at each step of the assignment.

## 1. Dataset Preparation

In order to prepare the data before preprocessing first I read columns "review\_body", and "star\_rating" from the CSV file to store as my data frame(as per the HW guidelines). After this, I separated my data into three classes (1, 2, 3) using a dictionary and added randomly shuffled the data to get different values at each execution. Finally, I picked 20000 rows from each of the three classes to create 60000 data set for further analysis.

*Average reviews length before data cleaning: 291.9329333333333*

## 2. Data Cleaning

For data cleaning, I have created a common function that will clean data for the following cases:

- Convert the reviews data into lowercase characters
- Remove numerical characters from the reviews data
- Remove punctuation marks from the reviews data
- Remove extra spaces from the reviews data
- Remove URLs from the reviews data
- Remove HTML tags from the reviews data

After this, I performed contradictions on the review data using a contradiction dictionary to further clean the data.

*Average reviews length after data cleaning: 280.53248333333335*

## 3. Data Preprocessing

For data preprocessing, I have used the NLTK package to first remove all the stop words from the dataset. After this, I performed lemmatization using WordNetLemmatizer from the NLTK package.

*Average review length after data preprocessing: 174.3622*

## 4. Feature Extraction

After data processing, I used TfidfVectorizer from sklearn package to extract TF-IDF features. I further divided the dataset into 80% training dataset and 20% testing dataset.

## 5. Results Of The Perceptron Model

	Precision	Recall	f1-score
Class 1	0.6309553819006806	0.6103389417215314	0.6204759543877045
Class 2	0.4973887092762994	0.5073566717402334	0.5023232450081627
Class 3	0.6622632103688934	0.671468284053576	0.6668339816790061
Average	0.5974024863809645	0.5966666666666667	0.5969493488558317

## 6. Results Of The SVM Model

	Precision	Recall	f1-score
Class 1	0.6896899420216789	0.6689486552567238	0.6791609780315253
Class 2	0.5369311116637653	0.5825688073394495	0.5588197230490488
Class 3	0.7592223330009971	0.72454804947668886	0.7414800389483934
Average	0.6668724375525644	0.66175	0.66382803145898

## 7. Results Of The Logistic Regression Model

	Precision	Recall	f1-score
Class 1	0.7070834383665239	0.6853163938431468	0.6960297766749379
Class 2	0.5792091519522506	0.6014979338842975	0.5901431648295958
Class 3	0.7567298105682951	0.7524163568773234	0.7545669193488257
Average	0.6825162612696973	0.6808333333333333	0.6656925870710259

## 8. Results Of The Naive Bayes Model

	Precision	Recall	f1-score
Class 1	0.6637257373329972	0.6941734774584761	0.6786082474226804
Class 2	0.6187515543397165	0.5682960255824577	0.5924514823193238
Class 3	0.7198404785643071	0.7542439279185166	0.7366407346001785
Average	0.6652229349188391	0.6674166666666667	0.6815468108102689

# NLP\_HW11

January 25, 2023

Importing All Required Libraries

```
[1]: import pandas as pd
import numpy as np
import random
from tqdm import tqdm
import nltk
```

```
[2]: cd /content/drive/MyDrive/NLP
```

/content/drive/MyDrive/NLP

```
[ ]: !wget https://s3.amazonaws.com/amazon-reviews-pds/tsv/
      ↪amazon_reviews_us_Beauty_v1_00.tsv.gz
```

```
--2023-01-25 03:40:38-- https://s3.amazonaws.com/amazon-reviews-
pds/tsv/amazon_reviews_us_Beauty_v1_00.tsv.gz
Resolving s3.amazonaws.com (s3.amazonaws.com)... 52.216.62.40, 52.217.132.0,
52.216.39.72, ...
Connecting to s3.amazonaws.com (s3.amazonaws.com)|52.216.62.40|:443...
connected.
HTTP request sent, awaiting response... 200 OK
Length: 914070021 (872M) [application/x-gzip]
Saving to: 'amazon_reviews_us_Beauty_v1_00.tsv.gz.1'
```

```
amazon_reviews_us_B 100%[=====>] 871.72M 54.1MB/s in 22s
```

```
2023-01-25 03:41:00 (40.1 MB/s) - 'amazon_reviews_us_Beauty_v1_00.tsv.gz.1'
saved [914070021/914070021]
```

Unzipping Reviews Data

```
[ ]: !gunzip amazon_reviews_us_Beauty_v1_00.tsv.gz
```

```
gzip: amazon_reviews_us_Beauty_v1_00.tsv already exists; do you wish to
overwrite (y or n)? y
```

Reading Data From CSV File

```
[3]: dataframe = pd.read_csv(r"amazon_reviews_us_Beauty_v1_00.tsv",sep="\t",usecols_
    ↪ ["review_body","star_rating"])
```

/usr/local/lib/python3.8/dist-packages/IPython/core/interactiveshell.py:3326:  
DtypeWarning: Columns (7) have mixed types.Specify dtype option on import or set  
low\_memory=False.

```
    exec(code_obj, self.user_global_ns, self.user_ns)
```

Separating Ratings Into Three Classes

```
[4]: ratingDict = {1:1, 2:1, 3:2, 4:3, 5:3}
    filteredDataFrame = dataframe.replace({"star_rating": ratingDict})
    filteredDataFrame = filteredDataFrame.sample(frac=1, random_state=14).
    ↪reset_index(drop=True)
    filteredDataFrame
```

```
[4]:      star_rating      review_body
0           3  There aren't nearly 600 pieces, but for the pr...
1           3  good stuff, good price - must use with its own...
2           2  I like the feel of the lipstick, but it's too ...
3           3  Ive been using this system fairly regularly fo...
4           3  I had a Remington makeup mirror for about a de...
...         ...
5094558      2          Dont work for my stinky butt
5094559      3  I just Love this colors! They is pretty and tr...
5094560      3  these nail art bows are beautiful bows n pearl...
5094561      5  Leaves my tangled messy long hair soft and smo...
5094562      5  I LOVE this cream!! Makes my hands and entire ...
```

[5094563 rows x 2 columns]

Contradictions Dictionary For Performing Contradictions On Reviews

```
[5]: contractions_dict = {
    "ain't": "are not",
    "aren't": "are not",
    "can't": "cannot",
    "can't've": "cannot have",
    "'cause": "because",
    "could've": "could have",
    "couldn't": "could not",
    "couldn't've": "could not have",
    "didn't": "did not",
    "doesn't": "does not",
    "don't": "do not",
    "hadn't": "had not",
    "hadn't've": "had not have",
    "hasn't": "has not",
```

"haven't": "have not",  
"he'd": "he would",  
"he'd've": "he would have",  
"he'll": "he will",  
"he'll've": "he will have",  
"he's": "he is",  
"how'd": "how did",  
"how'd'y": "how do you",  
"how'll": "how will",  
"how's": "how is",  
"I'd": "I had",  
"I'd've": "I would have",  
"I'll": "I will",  
"I'll've": "I will have",  
"I'm": "I am",  
"I've": "I have",  
"isn't": "is not",  
"it'd": "it had",  
"it'd've": "it would have",  
"it'll": "it will",  
"it'll've": "it will have",  
"it's": "it is",  
"let's": "let us",  
"ma'am": "madam",  
"mayn't": "may not",  
"might've": "might have",  
"mightn't": "might not",  
"mightn't've": "might not have",  
"must've": "must have",  
"mustn't": "must not",  
"mustn't've": "must not have",  
"needn't": "need not",  
"needn't've": "need not have",  
"o'clock": "of the clock",  
"oughtn't": "ought not",  
"oughtn't've": "ought not have",  
"shan't": "shall not",  
"sha'n't": "shall not",  
"shan't've": "shall not have",  
"she'd": "she would",  
"she'd've": "she would have",  
"she'll": "she will",  
"she'll've": "she will have",  
"she's": "she is",  
"should've": "should have",  
"shouldn't": "should not",  
"shouldn't've": "should not have",

"so've": "so have",  
"so's": "so is",  
"that'd": "that would",  
"that'd've": "that would have",  
"that's": "that is",  
"there'd": "there would",  
"there'd've": "there would have",  
"there's": "there is",  
"they'd": "they had",  
"they'd've": "they would have",  
"they'll": "they will",  
"they'll've": "they will have",  
"they're": "they are",  
"they've": "they have",  
"to've": "to have",  
"wasn't": "was not",  
"we'd": "we would",  
"we'd've": "we would have",  
"we'll": "we will",  
"we'll've": "we will have",  
"we're": "we are",  
"we've": "we have",  
"weren't": "were not",  
"what'll": "what will",  
"what'll've": "what will have",  
"what're": "what are",  
"what's": "what is",  
"what've": "what have",  
"when's": "when is",  
"when've": "when have",  
"where'd": "where did",  
"where's": "where is",  
"where've": "where have",  
"who'll": "who will",  
"who'll've": "who will have",  
"who's": "who is",  
"who've": "who have",  
"why's": "why is",  
"why've": "why have",  
"will've": "will have",  
"won't": "will not",  
"won't've": "will not have",  
"would've": "would have",  
"wouldn't": "would not",  
"wouldn't've": "would not have",  
"y'all": "you all",  
"y'all'd": "you all would",

```

"y'all'd've": "you all would have",
"y'all're": "you all are",
"y'all've": "you all have",
"you'd": "you would",
"you'd've": "you would have",
"you'll": "you will",
"you'll've": "you will have",
"you're": "you are",
"you've": "you have",
"i'd": "i would",
"i'd've": "i would have",
"i'll": "i will",
"i'll've": "i will have",
"i've": "i have"
}

```

Taking 20000 Data From Each Rating Class

```

[6]: rating1Data = filteredDataFrame[filteredDataFrame['star_rating']==1].head(20000)
rating2Data = filteredDataFrame[filteredDataFrame['star_rating']==2].head(20000)
rating3Data = filteredDataFrame[filteredDataFrame['star_rating']==3].head(20000)

finalRatingsData = rating1Data.append(rating2Data).append(rating3Data).
    ↪reset_index()
finalRatingsData

```

```

[6]:      index star_rating      review_body
0         6           1  While this cap fits really well, it smells hor...
1        18           1  I do not like it no is do not do nothing do no...
2        24           1  Bought this product a few months ago. Not happ...
3        37           1  Simple...right? Wrong. It's simple if it actua...
4        49           1  This candle has a nice container, lid, etc but...
...     ...           ...
59995  31882           3  I purchased this product because it is suppose...
59996  31885           3  It arrived all in one piece and it smells grea...
59997  31886           3  it goes on smoothly, spreads well and you do n...
59998  31888           3  Love opi products!!!!
59999  31889           3  Works great! Makes my beard shine (my wife say...

```

[60000 rows x 3 columns]

```

[7]: finalRatingsData['review_body'] = finalRatingsData["review_body"].apply(str)
    ↪#converts review_body column to string type

```

```

[8]: averageStringLengthBeforeDataCleaning, stringLength = 0, 0
for ratings in finalRatingsData['review_body']:
    stringLength += len(ratings)

```

```
averageStringLengthBeforeDataCleaning = stringLength /
↳len(finalRatingsData['review_body'])
print("Average reviews length before data cleaning : ",
↳averageStringLengthBeforeDataCleaning)
```

Average reviews length before data cleaning : 291.93293333333333

Performing Preprocessing/Cleaning Of Review Data

```
[9]: import re

def cleanReviewData(column):
    column = column.lower() #converts string to lowercase
    column = re.sub(r'\d+', '', column) #remove numerical characters from the string
    column = re.sub(r'[\W\s]', '', column) #removes punctuations from the string
    column = column.strip() #remove extra spaces from the string
    column = re.sub(r"http\S+", "", column) #removes URLs from the string
    column = re.sub(re.compile('<.*?>'), '', column) #removes HTML tags from the
↳string

    return column

finalRatingsData['review_body'] = finalRatingsData["review_body"].
↳apply(cleanReviewData)
finalRatingsData
```

```
[9]:      index star_rating      review_body
0         6          1  while this cap fits really well it smells horr...
1        18          1  i do not like it no is do not do nothing do no...
2        24          1  bought this product a few months ago not happy...
3        37          1  simpleright wrong its simple if it actually wo...
4        49          1  this candle has a nice container lid etc but h...
...     ...          ...
59995  31882          3  i purchased this product because it is suppose...
59996  31885          3  it arrived all in one piece and it smells grea...
59997  31886          3  it goes on smoothly spreads well and you do no...
59998  31888          3                      love opi products
59999  31889          3  works great makes my beard shine my wife says ...
```

[60000 rows x 3 columns]

Using Contradictions Dictionary To Perform Contradictions

```
[10]: for reviews in finalRatingsData['review_body']:
        review = reviews.split()
        for char in review:
```



```

    if char in contractions_dict:
        reviews = reviews.replace(char, contractions_dict[char])

finalRatingsData

```

```

[10]:      index star_rating      review_body
0         6          1  while this cap fits really well it smells horr...
1        18          1  i do not like it no is do not do nothing do no...
2        24          1  bought this product a few months ago not happy...
3        37          1  simplerright wrong its simple if it actually wo...
4        49          1  this candle has a nice container lid etc but h...
...      ...      ...
59995  31882          3  i purchased this product because it is suppose...
59996  31885          3  it arrived all in one piece and it smells grea...
59997  31886          3  it goes on smoothly spreads well and you do no...
59998  31888          3                      love opi products
59999  31889          3  works great makes my beard shine my wife says ...

```

[60000 rows x 3 columns]

```

[11]: averageStringLengAfterDataCleaning, stringLength = 0, 0
      for ratings in finalRatingsData['review_body']:
          stringLength += len(ratings)

      averageStringLengAfterDataCleaning = stringLength /
      ↪len(finalRatingsData['review_body'])
      print("Average reviews length after data cleaning : ",
      ↪averageStringLengAfterDataCleaning)

```

Average reviews length after data cleaning : 280.53248333333335

```

[12]: averageStringLengBeforeDataPreprocessing, stringLength = 0, 0
      for ratings in finalRatingsData['review_body']:
          stringLength += len(ratings)

      averageStringLengBeforeDataPreprocessing = stringLength /
      ↪len(finalRatingsData['review_body'])
      print("Average reviews length before data preprocessing : ",
      ↪averageStringLengBeforeDataPreprocessing)

```

Average reviews length before data preprocessing : 280.53248333333335

Using NLTK Library To Remove Stop Words

```

[13]: nltk.download('stopwords')
      from nltk.corpus import stopwords
      stop_words = set(stopwords.words('english'))

```

```
finalRatingsData['review_body'] = finalRatingsData['review_body'].apply(lambda_
↪key: ' '.join([word for word in key.split() if word not in (stop_words)]))
```

[nltk\_data] Downloading package stopwords to /root/nltk\_data...

[nltk\_data] Package stopwords is already up-to-date!

Using NLTK Library To Perform Lemmatization

```
[14]: from nltk.stem import WordNetLemmatizer
nltk.download('wordnet')
nltk.download('omw-1.4')
w_tokenizer = nltk.tokenize.WhitespaceTokenizer()
lemmatizer = nltk.stem.WordNetLemmatizer()
def lemmatize_text(text):
    return " ".join([lemmatizer.lemmatize(w) for w in w_tokenizer.
↪tokenize(text)])
finalRatingsData['review_body'] = finalRatingsData['review_body'].
↪apply(lemmatize_text)
```

[nltk\_data] Downloading package wordnet to /root/nltk\_data...

[nltk\_data] Package wordnet is already up-to-date!

[nltk\_data] Downloading package omw-1.4 to /root/nltk\_data...

[nltk\_data] Package omw-1.4 is already up-to-date!

```
[15]: averageStringLengAfterDataPreprocessing, stringLength = 0, 0
for ratings in finalRatingsData['review_body'].to_list():
    stringLength += len(ratings)

averageStringLengAfterDataPreprocessing = stringLength /_
↪len(finalRatingsData['review_body'])
print("Average reviews length after data preprocessing : ",_
↪averageStringLengAfterDataPreprocessing)
```

Average reviews length after data preprocessing : 174.3622

Feature Extraction Using TF-IDF

```
[16]: print(finalRatingsData)
from sklearn.feature_extraction.text import TfidfVectorizer
tfidfvectorizer = TfidfVectorizer()
x = tfidfvectorizer.fit_transform(finalRatingsData['review_body'])
x
```

	index	star_rating	review_body
0	6	1	cap fit really well smell horrible tried washi...
1	18	1	like nothing west money personal opinion sorry
2	24	1	bought product month ago happy result image wo...
3	37	1	simpleright wrong simple actually work sorry d...
4	49	1	candle nice container lid etc scent little bur...

```

...      ...      ...
59995  31882      3  purchased product supposed help hair loss im sure
59996  31885      3  arrived one piece smell great like always husb..
59997  31886      3  go smoothly spread well need much cover face n..
59998  31888      3                               love opi product
59999  31889      3    work great make beard shine wife say smell good

```

[60000 rows x 3 columns]

```

[16]: <60000x47252 sparse matrix of type '<class 'numpy.float64'>'
      with 1374165 stored elements in Compressed Sparse Row format>

```

Splitting Data Into Training And Testing Set

```

[17]: from sklearn.model_selection import train_test_split
      part1 = finalRatingsData['review_body'].to_list()
      part2 = finalRatingsData['star_rating'].to_list()
      part1_train, part1_test, part2_train, part2_test = train_test_split(part1,
      ↪part2, test_size=0.2, shuffle = True)

```

```

[18]: Train_X_Tfidf = tfidfvectorizer.transform(part1_train)
      Test_X_Tfidf = tfidfvectorizer.transform(part1_test)

```

Running Naive Bayes Model

```

[19]: from sklearn import model_selection, naive_bayes, svm
      from sklearn.metrics import accuracy_score

      Naive = naive_bayes.MultinomialNB()
      Naive.fit(Train_X_Tfidf,part2_train)
      predictions_NB = Naive.predict(Test_X_Tfidf)
      print("Naive Bayes Accuracy Score -> ",accuracy_score(predictions_NB,
      ↪part2_test)*100)

```

Naive Bayes Accuracy Score -> 66.74166666666666

Running SVM Model

```

[20]: from sklearn.svm import LinearSVC

      svm = LinearSVC()
      svm.fit(Train_X_Tfidf,part2_train)
      predictions_svm = svm.predict(Test_X_Tfidf)
      print("SVM Accuracy Score -> ",accuracy_score(predictions_svm, part2_test)*100)

```

SVM Accuracy Score -> 66.175

Running Perceptron Model

```
[21]: from sklearn.linear_model import Perceptron

perceptron = Perceptron()
perceptron.fit(Train_X_Tfidf,part2_train)
predictions_perp = perceptron.predict(Test_X_Tfidf)
print("Perceptron Accuracy Score -> ",accuracy_score(predictions_perp,
↪part2_test)*100)
```

Perceptron Accuracy Score -> 59.66666666666667

Running Logistic Regression Model

```
[22]: from sklearn.linear_model import LogisticRegression

logisticRegression = LogisticRegression(max_iter=1000)
logisticRegression.fit(Train_X_Tfidf,part2_train)
predictions_log = logisticRegression.predict(Test_X_Tfidf)
print("LogisticRegression Accuracy Score -> ",accuracy_score(predictions_log,
↪part2_test)*100)
```

LogisticRegression Accuracy Score -> 68.08333333333333

```
[23]: from sklearn import metrics

predictions_NB_output = metrics.classification_report(predictions_NB,
↪part2_test, output_dict=True)
print("Naive Bayes Output -> ",predictions_NB_output)
```

Naive Bayes Output -> {'1': {'precision': 0.6637257373329972, 'recall': 0.6941734774584761, 'f1-score': 0.6786082474226804, 'support': 3793}, '2': {'precision': 0.6187515543397165, 'recall': 0.5682960255824577, 'f1-score': 0.5924514823193238, 'support': 4378}, '3': {'precision': 0.7198404785643071, 'recall': 0.7542439279185166, 'f1-score': 0.7366407346001785, 'support': 3829}, 'accuracy': 0.6674166666666667, 'macro avg': {'precision': 0.6674392567456735, 'recall': 0.6722378103198169, 'f1-score': 0.6692334881140609, 'support': 12000}, 'weighted avg': {'precision': 0.6652229349188391, 'recall': 0.6674166666666667, 'f1-score': 0.6656925870710259, 'support': 12000}}

Output Of Naive Bayes Model

```
[24]: for key, val in predictions_NB_output.items():
    if key == "accuracy" or key == "macro avg":
        continue
    if key == "1":
        print("Class 1 Precison : ", val['precision'])
        print("Class 1 Recall : ", val['recall'])
        print("Class 1 f1-score : ", val['f1-score'])
    elif key == "2":
        print("Class 2 Precison : ", val['precision'])
```

```

        print("Class 2 Recall : ", val['recall'])
        print("Class 2 f1-score : ", val['f1-score'])
    elif key == "3":
        print("Class 3 Precison : ", val['precision'])
        print("Class 3 Recall : ", val['recall'])
        print("Class 3 f1-score : ", val['f1-score'])
    elif key == "weighted avg":
        print("Average Precison : ", val['precision'])
        print("Average Recall : ", val['recall'])
        print("Average f1-score : ", val['f1-score'])

```

```

Class 1 Precison : 0.6637257373329972
Class 1 Recall : 0.6941734774584761
Class 1 f1-score : 0.6786082474226804
Class 2 Precison : 0.6187515543397165
Class 2 Recall : 0.5682960255824577
Class 2 f1-score : 0.5924514823193238
Class 3 Precison : 0.7198404785643071
Class 3 Recall : 0.7542439279185166
Class 3 f1-score : 0.7366407346001785
Average Precison : 0.6652229349188391
Average Recall : 0.6674166666666667
Average f1-score : 0.6656925870710259

```

[25]: `from sklearn import metrics`

```

predictions_svm_output = metrics.classification_report(predictions_svm,
↳part2_test, output_dict=True)
print("SVM Output -> ", predictions_svm_output)

```

```

SVM Output -> {'1': {'precision': 0.6896899420216789, 'recall':
0.6689486552567238, 'f1-score': 0.6791609780315253, 'support': 4090}, '2':
{'precision': 0.5369311116637653, 'recall': 0.5825688073394495, 'f1-score':
0.5588197230490488, 'support': 3706}, '3': {'precision': 0.7592223330009971,
'recall': 0.7245480494766888, 'f1-score': 0.7414800389483934, 'support': 4204},
'accuracy': 0.66175, 'macro avg': {'precision': 0.6619477955621471, 'recall':
0.6586885040242874, 'f1-score': 0.6598202466763224, 'support': 12000}, 'weighted
avg': {'precision': 0.6668724375525644, 'recall': 0.66175, 'f1-score':
0.66382803145898, 'support': 12000}}

```

Output Of SVM Model

[26]:

```

for key, val in predictions_svm_output.items():
    if key == "accuracy" or key == "macro avg":
        continue
    if key == "1":
        print("Class 1 Precison : ", val['precision'])
        print("Class 1 Recall : ", val['recall'])

```

```

        print("Class 1 f1-score : ", val['f1-score'])
    elif key == "2":
        print("Class 2 Precison : ", val['precision'])
        print("Class 2 Recall : ", val['recall'])
        print("Class 2 f1-score : ", val['f1-score'])
    elif key == "3":
        print("Class 3 Precison : ", val['precision'])
        print("Class 3 Recall : ", val['recall'])
        print("Class 3 f1-score : ", val['f1-score'])
    elif key == "weighted avg":
        print("Average Precison : ", val['precision'])
        print("Average Recall : ", val['recall'])
        print("Average f1-score : ", val['f1-score'])

```

```

Class 1 Precison : 0.6896899420216789
Class 1 Recall : 0.6689486552567238
Class 1 f1-score : 0.6791609780315253
Class 2 Precison : 0.5369311116637653
Class 2 Recall : 0.5825688073394495
Class 2 f1-score : 0.5588197230490488
Class 3 Precison : 0.7592223330009971
Class 3 Recall : 0.7245480494766888
Class 3 f1-score : 0.7414800389483934
Average Precison : 0.6668724375525644
Average Recall : 0.66175
Average f1-score : 0.66382803145898

```

```

[27]: from sklearn import metrics

predictions_perp_output = metrics.classification_report(predictions_perp,
↳ part2_test, output_dict=True)
print("Perceptron Output -> ", predictions_perp_output)

```

```

Perceptron Output -> {'1': {'precision': 0.6309553819006806, 'recall':
0.6103389417215314, 'f1-score': 0.6204759543877045, 'support': 4101}, '2':
{'precision': 0.4973887092762994, 'recall': 0.5073566717402334, 'f1-score':
0.5023232450081627, 'support': 3942}, '3': {'precision': 0.6622632103688934,
'recall': 0.671468284053576, 'f1-score': 0.6668339816790061, 'support': 3957},
'accuracy': 0.5966666666666667, 'macro avg': {'precision': 0.5968691005152911,
'recall': 0.5963879658384469, 'f1-score': 0.5965443936916245, 'support': 12000},
'weighted avg': {'precision': 0.5974024863809645, 'recall': 0.5966666666666667,
'f1-score': 0.5969493488558317, 'support': 12000}}

```

Output Of Perceptron Model

```

[28]: for key, val in predictions_perp_output.items():
        if key == "accuracy" or key == "macro avg":
            continue

```

```

if key == "1":
    print("Class 1 Precison : ", val['precision'])
    print("Class 1 Recall : ", val['recall'])
    print("Class 1 f1-score : ", val['f1-score'])
elif key == "2":
    print("Class 2 Precison : ", val['precision'])
    print("Class 2 Recall : ", val['recall'])
    print("Class 2 f1-score : ", val['f1-score'])
elif key == "3":
    print("Class 3 Precison : ", val['precision'])
    print("Class 3 Recall : ", val['recall'])
    print("Class 3 f1-score : ", val['f1-score'])
elif key == "weighted avg":
    print("Average Precison : ", val['precision'])
    print("Average Recall : ", val['recall'])
    print("Average f1-score : ", val['f1-score'])

```

```

Class 1 Precison : 0.6309553819006806
Class 1 Recall : 0.6103389417215314
Class 1 f1-score : 0.6204759543877045
Class 2 Precison : 0.4973887092762994
Class 2 Recall : 0.5073566717402334
Class 2 f1-score : 0.5023232450081627
Class 3 Precison : 0.6622632103688934
Class 3 Recall : 0.671468284053576
Class 3 f1-score : 0.6668339816790061
Average Precison : 0.5974024863809645
Average Recall : 0.5966666666666667
Average f1-score : 0.5969493488558317

```

[29]: `from sklearn import metrics`

```

predictions_log_output = metrics.classification_report(predictions_log,
↳ part2_test, output_dict=True)
print("LogisticRegression Output -> ", predictions_log_output)

```

```

LogisticRegression Output -> {'1': {'precision': 0.7070834383665239, 'recall':
0.6853163938431468, 'f1-score': 0.6960297766749379, 'support': 4093}, '2':
{'precision': 0.5792091519522506, 'recall': 0.6014979338842975, 'f1-score':
0.5901431648295958, 'support': 3872}, '3': {'precision': 0.7567298105682951,
'recall': 0.7524163568773234, 'f1-score': 0.7545669193488257, 'support': 4035},
'accuracy': 0.6808333333333333, 'macro avg': {'precision': 0.6810074669623566,
'recall': 0.6797435615349227, 'f1-score': 0.6802466202844532, 'support': 12000},
'weighted avg': {'precision': 0.6825162612696973, 'recall': 0.6808333333333333,
'f1-score': 0.6815468108102689, 'support': 12000}}

```

Output Of Logistic Regression Model

```
[30]: for key, val in predictions_log_output.items():
    if key == "accuracy" or key == "macro avg":
        continue
    if key == "1":
        print("Class 1 Precison : ", val['precision'])
        print("Class 1 Recall : ", val['recall'])
        print("Class 1 f1-score : ", val['f1-score'])
    elif key == "2":
        print("Class 2 Precison : ", val['precision'])
        print("Class 2 Recall : ", val['recall'])
        print("Class 2 f1-score : ", val['f1-score'])
    elif key == "3":
        print("Class 3 Precison : ", val['precision'])
        print("Class 3 Recall : ", val['recall'])
        print("Class 3 f1-score : ", val['f1-score'])
    elif key == "weighted avg":
        print("Average Precison : ", val['precision'])
        print("Average Recall : ", val['recall'])
        print("Average f1-score : ", val['f1-score'])
```

```
Class 1 Precison : 0.7070834383665239
Class 1 Recall : 0.6853163938431468
Class 1 f1-score : 0.6960297766749379
Class 2 Precison : 0.5792091519522506
Class 2 Recall : 0.6014979338842975
Class 2 f1-score : 0.5901431648295958
Class 3 Precison : 0.7567298105682951
Class 3 Recall : 0.7524163568773234
Class 3 f1-score : 0.7545669193488257
Average Precison : 0.6825162612696973
Average Recall : 0.6808333333333333
Average f1-score : 0.6815468108102689
```